


On-board fish detector using machine learning


Clare Walker, August 2021

Aim: Automate fish detection for passive imaging systems



FishOASIS imaging and acoustic system deployed in kelp forest MPAs

- 1500 photos a day stored on R-Pi
- All fish counted by scientists



Incorporate on-board ML tool to automate detection and counting

- Only store photos with fish
- Only check sample of labels

Lit review

Observations from literature review

- Much of existing literature applying CNNs to images of fish focus on classification
- Object detection has predominantly been used for post-processing in labs
- Open-source tools exist for this context: VIAME, FathomNet, Fish4Knowledge
- Results confirm that YOLO is faster than region proposal techniques, but requires more training data (>1000 annotations of each class)
- Most results come from coral reefs and other shallow, relatively static environments – no evidence (so far) of application in kelp forests
- Tools for on-board, real-time detection have been developed for midwater ROVs rather than passive, low-cost imaging systems

Most similar paper by Coro, to be published July 2021

Similarities:

- On-board object detection for underwater applications
- Built as an 'intelligent' component of passive monitoring system
- Open-source code, deployable after specifying a few parameters
- Utilizes YOLO Tiny

Differences:

- Optimised for large fish, rather than detecting all fish in an image
- Tested in relatively static environments, rather than in dynamic kelp forests
- No re-training for use in new environments, harder to update based on experience
- Training set generated by animating fish, rather than on live data
- Utilizes YOLO Tiny v3 rather than v4, as well as options for two other models

Data & Methodology

Overview of key steps followed

1. Trained YOLO Tiny v4 on four different training sets:
 - a. Excluded images before / after cut-off times (determined by brightness)
 - b. Excluded images with no fish, i.e. downsampled majority class
 - c. Incorporated simple batch image enhancements (edge enhance using PIL)
 - d. Removed labels for small fish below YOLO recommended size (see analysis of results)
2. Calculated performance metrics to reflect keep/discard precision, recall, F1
3. Apply trained model to test sets
4. Create data visualisations for results (e.g. keep/discard heatmap, count time series)
5. Clean up notebooks and create GitHub repo

YOLO v4 Tiny architecture

layer	filters	size/strd(dil)	input	output
0 conv	32	3 x 3/ 2	608 x 608 x 3 ->	304 x 304 x 32 0.160 BF
1 conv	64	3 x 3/ 2	304 x 304 x 32 ->	152 x 152 x 64 0.852 BF
2 conv	64	3 x 3/ 1	152 x 152 x 64 ->	152 x 152 x 64 1.703 BF
3 route	2		1/2 ->	152 x 152 x 32
4 conv	32	3 x 3/ 1	152 x 152 x 32 ->	152 x 152 x 32 0.426 BF
5 conv	32	3 x 3/ 1	152 x 152 x 32 ->	152 x 152 x 32 0.426 BF
6 route	5 4		->	152 x 152 x 64
7 conv	64	1 x 1/ 1	152 x 152 x 64 ->	152 x 152 x 64 0.189 BF
8 route	2 7		->	152 x 152 x 128
9 max		2x 2/ 2	152 x 152 x 128 ->	76 x 76 x 128 0.003 BF
10 conv	128	3 x 3/ 1	76 x 76 x 128 ->	76 x 76 x 128 1.703 BF
11 route	10		1/2 ->	76 x 76 x 64
12 conv	64	3 x 3/ 1	76 x 76 x 64 ->	76 x 76 x 64 0.426 BF
13 conv	64	3 x 3/ 1	76 x 76 x 64 ->	76 x 76 x 64 0.426 BF
14 route	13 12		->	76 x 76 x 128
15 conv	128	1 x 1/ 1	76 x 76 x 128 ->	76 x 76 x 128 0.189 BF
16 route	10 15		->	76 x 76 x 256
17 max		2x 2/ 2	76 x 76 x 256 ->	38 x 38 x 256 0.001 BF
18 conv	256	3 x 3/ 1	38 x 38 x 256 ->	38 x 38 x 256 1.703 BF
19 route	18		1/2 ->	38 x 38 x 128
20 conv	128	3 x 3/ 1	38 x 38 x 128 ->	38 x 38 x 128 0.426 BF
21 conv	128	3 x 3/ 1	38 x 38 x 128 ->	38 x 38 x 128 0.426 BF
22 route	21 20		->	38 x 38 x 256
23 conv	256	1 x 1/ 1	38 x 38 x 256 ->	38 x 38 x 256 0.189 BF
24 route	18 23		->	38 x 38 x 512
25 max		2x 2/ 2	38 x 38 x 512 ->	19 x 19 x 512 0.001 BF
26 conv	512	3 x 3/ 1	19 x 19 x 512 ->	19 x 19 x 512 1.703 BF
27 conv	256	1 x 1/ 1	19 x 19 x 512 ->	19 x 19 x 256 0.095 BF
28 conv	512	3 x 3/ 1	19 x 19 x 256 ->	19 x 19 x 512 0.852 BF
29 conv	18	1 x 1/ 1	19 x 19 x 512 ->	19 x 19 x 18 0.007 BF
30 yolo				

[yolo] params: iou_loss: ciou (4), iou_norm: 0.07, obj_norm: 1.00, cls_norm: 1.00, delta_norm: 1.00, scale_x_y: 1.05
nms_kind: greedy_nms (1), beta = 0.600000

31 route	27		->	19 x 19 x 256
32 conv	128	1 x 1/ 1	19 x 19 x 256 ->	19 x 19 x 128 0.024 BF
33 upsample		2x	19 x 19 x 128 ->	38 x 38 x 128
34 route	33 23		->	38 x 38 x 384
35 conv	256	3 x 3/ 1	38 x 38 x 384 ->	38 x 38 x 256 2.555 BF
36 conv	18	1 x 1/ 1	38 x 38 x 256 ->	38 x 38 x 18 0.013 BF
37 yolo				

[yolo] params: iou_loss: ciou (4), iou_norm: 0.07, obj_norm: 1.00, cls_norm: 1.00, delta_norm: 1.00, scale_x_y: 1.05
nms_kind: greedy_nms (1), beta = 0.600000

More information on object detection using CNNs and further explanation of YOLO found [here](#)

Summary of data split

Date	Site C	Site C.2			
		1	2	3	4
25/05/18					
26/05/18	5				
10/07/18					
11/07/18					
12/07/18					
13/07/18					4
14/07/18					
15/07/18					
16/07/18					
17/07/18		1	3		
18/07/18					
19/07/18					
20/07/18					
21/07/18		2	6		
22/07/18					
23/07/18					

- 1 **Training & test:** Same site, same perspective, same days, different images
- 2 **Test:** Same site, same perspective, different days (temporal)
- 3 **Test:** Same site, same days, different perspective (spatial)
- 4 **Test:** Same site, different days, different perspective (temporal and spatial)
- 5 **Test:** Different site, different days (temporal and spatial)
- 6 **Proof of application:** Apply to unlabeled data, compare time series

Definition of performance metrics

Problem:

- We want to measure performance terms of binary classification: Keep or Discard
- The test data are imbalanced, typically **20% keep (at least one fish)** and **80% discard (no fish)**

Best practice:

- **Majority class** is typically referred to as the **negative outcome** (e.g. “no change” or “negative test result”)
- **Minority class** is typically referred to as the **positive outcome** (e.g. “change” or “positive test result”)
- This is because metrics measure ability to detect positive outcomes — if majority class were defined as positive outcome, a model that only predicted the majority class would look better than random

Solution:

- We will define **Keep as positive outcome** and **Discard as negative outcome**
- Note that due to above, results would look much better if defined in opposite direction

Definition of performance metrics

Measure in terms of binary classification: **Keep (Positive) / Discard (Negative)**

- True positive (TP) = found > 0 fish, image had > 0 fish
- False positive (FP) = found > 0 fish, image had no fish (error - okay)
- True negative (TN) = found no fish, image had no fish
- False negative (FN) = found no fish, image had > 0 fish (error - not okay)

Key metrics:

- Precision = $TP / (TP + FP)$: proportion of images kept that actually contained fish
- Recall = $TP / (TP + FN)$: proportion of images kept out all images that should have been kept
- F1 = $2 * Pr * Re / (Pr + Re)$: harmonic mean of precision and recall

Example: Choice of positive outcome impacts results

Positive outcome: Keep

Confusion matrix

	Kept (+)	Discarded (-)	Total
Has fish (+)	133 (TP)	64 (FN)	197
Has no fish (-)	41 (FP)	762 (TN)	803
Total	174	826	1000

Accuracy = $895 / 1000 = 90\%$

Precision = $133 / 174 = 76\%$

Recall = $133 / 197 = 68\%$

Positive outcome: Discard

Confusion matrix

	Discarded (+)	Kept (-)	Total
Has no fish (+)	762 (TP)	41 (FN)	803
Has fish (-)	64 (FP)	133 (TN)	197
Total	826	174	1000

Accuracy = $895 / 1000 = 90\%$

Precision = $762 / 826 = 92\%$

Recall = $762 / 803 = 95\%$

Results

Preliminary results (imbalanced test sets)

		Model a ($n = 4171$, % pos = 20)			Model b ($n = 837$, % pos = 100)		
	% pos	Precision	Recall	F1	Precision	Recall	F1
Valid (1)	20	82	49	61	76	68	72
Valid (1) - no small fish	9 (a) 8 (b)	39	50	44	40	81	53
Test (2)	11	57	46	51	48	63	54
Test (3)	13	48	43	45	31	66	42
Test (4)	37	61	55	58	51	70	59
Test (5)	32	32	70	44	32	80	46

Preliminary results (balanced test sets) - *better precision*

		Model a ($n = 4171$, % pos = 20)			Model b ($n = 837$, % pos = 100)		
	% pos	Precision	Recall	F1	Precision	Recall	F1
Valid (1)	50	95	49	65	94	68	78
Valid (1) - no small fish	50	85	50	63	92	81	86
Test (2)	50	95	46	62	97	63	77
Test (3)	50	95	43	59	80	66	72
Test (4)	50	74	55	63	66	70	68
Test (5)	50	47	70	56	48	80	60

Analysis of results: Where is the model going wrong?

False negatives

This occurs when the model fails to detect fish in an image, which is usually owing to:

1. Uncommon/rare fish not in training set
2. Labelled fish are **blurry/unclear** *
3. Labelled fish are **too small** **

False positives

This occurs when the model confuses something else for a fish, usually:

1. Slight variation in water color
2. Kelp



Conducted two further experiments:



- c. Incorporated simple batch image enhancements from PIL (addresses *)
- d. Removed labels for small fish below YOLO recommended size (addresses **)

Analysis of results: Uncommon/rare fish (FN)






Annotated results sorted by TP, TN, FP, FN are saved [here](#)

Analysis of results: Labelled fish too blurry (FN)




Annotated results sorted by TP, TN, FP, FN are saved [here](#)

Analysis of results: Labelled fish too blurry (FP)





Annotated results sorted by TP, TN, FP, FN are saved [here](#)

Analysis of results: Labelled fish too small (FN)




Annotated results sorted by TP, TN, FP, FN are saved [here](#)

Analysis of results: Kelp confused with fish (FP)



Annotated results sorted by TP, TN, FP, FN are saved [here](#)

Example: Best image processing from PIL (Model c)



```
def enhance(path):  
    im = Image.open(path)  
    im = im.filter(EDGE_ENHANCE_MORE)  
    enh = ImageEnhance.Contrast(im)  
    return enh.enhance(1.2)
```

Preliminary results (imbalanced test sets)

		Model c ($n = 837$, % pos = 100)			Model d ($n = 837$, % pos = 47)		
	% pos	Precision	Recall	F1	Precision	Recall	F1
Valid (1)	20 / 8	59	46	52	69	44	53
Test (2)							
Test (3)							
Test (4)							
Test (5)							

Preliminary results (balanced test sets) - *better precision*

		Model c ($n = 837$, % pos = 100)			Model d ($n = 837$, % pos = 47)		
	% pos	Precision	Recall	F1	Precision	Recall	F1
Valid (1)	50	86	46	60	97	44	60
Test (2)							
Test (3)							
Test (4)							
Test (5)							

Conclusion

Conclusion: Further improvement limited by 5 challenges

Key takeaways:

- Best results achieved with **model b** (trained on only positive examples) and tested on set **without small fish**
- Given existing labels and basic image enhancement packages, this appears to be an **upper bound on performance**

Further improvements are limited by:

1. Many fish of interest are too small for basic YOLO v4 implementation
2. Some labels are too blurry to make out, especially without bespoke editing (e.g. in Adobe LightRoom)
3. Fish are too heterogenous to treat as one class
4. Real data is imbalanced (i.e. % images without fish > % images with fish)
5. Kelp can look like fish resulting in false positives (less severe)

These could be addressed by further research:

- Define and apply more sophisticated image processing algorithms (e.g. <https://colorcorrection.firebaseio.com/>)
- Use another version of YOLO optimised for small objects
- Relabel images with only distinct fish

However, there is no obvious solution for class imbalance, fish heterogeneity and false positives due to kelp


Colab Notebooks will be shared on GitHub (w/o data)


FishOASIS_ML-Detector > CLARE > Notebooks ▾


Name ↑

 Mat to YOLO conversion.ipynb

 YoloV4-FishOASIS-Test-Analyze.ipynb

 YoloV4-FishOASIS-Train-A.ipynb

 YoloV4-FishOASIS-Train-B.ipynb

 YoloV4-FishOASIS-Train-C.ipynb


 YoloV4-FishOASIS-Train-D.ipynb

Next steps


1. Try a few more experiments with YOLO Tiny v4:
 - a. “Clean” labels (remove unclassified) and increase input image size
 - b. “Cleaner” labels (remove unclassified, remove edge fish, remove small fish) and increase input image size
 - c. Split image into e.g. 6 segments
 - d. Apply Derya’s image enhancement algorithm + other literature, using color palette
2. Treat as pure classification task, i.e. try experiment with VGG or ResNet

Heatmaps


Heatmaps | Model a, Valid (1)




Heatmaps | Model b, Valid (1)




Heatmaps | Model c, Valid (1)




Heatmaps | Model d, Valid (1)




Heatmaps | Model a, Test (2)




Heatmaps | Model b, Test (2)




Heatmaps | Model a, Test (3)




Heatmaps | Model b, Test (3)




Heatmaps | Model a, Test (4)



Heatmaps | Model b, Test (4)



Heatmaps | Model a, Test (5)



Heatmaps | Model b, Test (5)

